

# Operating Systems Lab



## Lab # 12

### Threads

Instructor: Engr. Muhammad Usman

Email: [usman.rafiq@nu.edu.pk](mailto:usman.rafiq@nu.edu.pk)

Course Code: CL2006

Department of Computer Science,  
National University of Computer and Emerging Sciences FAST  
Peshawar Campus

## Contents

POSIX Threads	<b>3</b>
What is a Thread?	<b>3</b>
Why Multithreading?	<b>3</b>
Thread Basics	<b>3</b>
Thread Creation	<b>4</b>
Example # 01:	6
Example # 02:	8
Passing Multiple Arguments to Thread	<b>10</b>
Example # 03	11
Thread Termination	<b>13</b>
Example # 04	14
Data Sharing between Threads	<b>15</b>
Example # 05	15
Synchronization through Mutex	<b>17</b>
Example # 06	18

# POSIX Threads

## POSIX Threads

POSIX thread library is a standard thread library for C/C++. Using the POSIX thread library, we can create a new concurrent process execution flow such that our program can then handle multiple execution paths. As a result, our program would appear to do many things at the same time.

## What is a Thread?

A thread is a single sequence stream within in a process. Because threads have some of the properties of processes, they are sometimes called *lightweight processes*.

## What are the differences between process and thread?

Threads are not independent of one other like processes as a result threads shares with other threads their **code section**, **data section** and OS resources like open files and signals. But, like process, a thread has its own program counter (PC), a register set, and a stack space.

## Why Multithreading?

Threads are popular way to improve application through parallelism. For example, in a browser, multiple tabs can be different threads. MS word uses multiple threads, one thread to format the text, other thread to process inputs, etc.

Threads operate faster than processes due to following reasons:

- 1) Thread creation is much faster.
- 2) Context switching between threads is much faster.
- 3) Threads can be terminated easily
- 4) Communication between threads is faster.

## Thread Basics

A thread can be defined as a separate stream of instructions within a process. From developer point of view, a thread is simply a function or procedure, that has its own existence and runs independently from the program's `main()` procedure/function.

To visualize, imagine a program C program with a number of functions. This program can be run by entering the executable **a.out** on the command interface. Then imagine each function being scheduled by the operating system. This would be a **multi-threaded** program.

Unlike child processes, a thread doesn't know which thread is responsible for its creation, neither does it maintain a list of current active threads inside a process. Within the same process, a thread may share the process instructions (text section), global data (data section) and open files (file descriptors). Each thread has a unique **thread ID**, **set of registers** and **stack**. These will be individual to a thread itself.

## Thread Creation

A thread is created using the `pthread_create()` call. Just like process creation using `fork()`, a `pthread_create()` call will return certain integer numbers upon successful completion of the call.

In a **Unix/Linux operating system**, the **C/C++ languages** provide the POSIX thread (pthread) standard API (Application program Interface) for all thread related functions. It allows us to create multiple threads for concurrent process flow. It is most effective on multiprocessor or multi-core systems where threads can be implemented on a kernel-level for achieving the speed of execution. Gains can also be found in uni-processor systems by exploiting the latency in IO or other system functions that may halt a process.

We must include the **pthread.h** header file at the beginning of the script to use all the functions of the pthreads library.

The **functions** defined in the **pthread library** include:

1. **pthread\_create**: used to create a new thread  
**Syntax:**

```
int pthread_create(pthread_t * thread,
                  const pthread_attr_t * attr,
                  void * (*start_routine)(void *),
                  void *arg);
```

### Parameters:

**thread:** pointer to an unsigned integer value that returns the thread id of the thread created.

**attr:** pointer to a structure that is used to define thread attributes like detached state, scheduling policy, stack address, etc. Set to **NULL** for default thread attributes.

**start\_routine:** pointer to a subroutine that is executed by the thread. The return type and parameter type of the subroutine must be of type void \*. The function has a single attribute but if multiple values need to be passed to the function, a **struct** must be used.

**arg:** pointer to void that contains the arguments to the function defined in the earlier argument.

### 2. **pthread\_exit:** used to terminate a thread

#### Syntax:

```
void pthread_exit(void *retval);
```

**Parameters:** This method accepts a mandatory parameter **retval** which is the pointer to an integer that stores the return status of the thread terminated. The scope of this variable must be global so that any thread waiting to join this thread may read the return status.

### 3. **pthread\_join:** used to wait for the termination of a thread.

#### Syntax:

```
int pthread_join(pthread_t th,
                 void **thread_return);
```

**Parameter:** This method accepts following parameters:

**th:** thread id of the thread for which the current thread waits.

**thread\_return:** pointer to the location where the exit status of the thread mentioned in th is stored.

4. **pthread\_self**: used to get the thread id of the current thread.

**Syntax:**

```
pthread_t pthread_self(void);
```

5. **pthread\_equal**: compares whether two threads are the same or not. If the two threads are equal, the function returns a non-zero value otherwise zero.

**Syntax:**

```
int pthread_equal(pthread_t t1,  
                  pthread_t t2);
```

**Parameters:** This method accepts following parameters:

- a. t1: the thread id of the first thread
- b. t2: the thread id of the second thread

## Example # 01:

### A simple C program to demonstrate use of pthread basic functions

Please note that the below program may compile only with C compilers with pthread library.

#### thread.c

```
#include <stdio.h>  
#include <stdlib.h>  
#include <unistd.h>  
#include <pthread.h>  
// A normal C function that is executed as a thread  
// when its name is specified in pthread_create()  
void *myThreadFun()  
{  
    sleep(1);  
    printf("Printing Hello World from Thread \n");  
}  
int main()  
{  
    pthread_t thread_id;  
    printf("Before Thread\n");
```

```
pthread_create(&thread_id, NULL, myThreadFun, NULL);  
pthread_join(thread_id, NULL);  
printf("After Thread\n");  
exit(0); }
```

In main() we declare a variable called thread\_id, which is of type pthread\_t, which is an integer used to identify the thread in the system. After declaring thread\_id, we call pthread\_create() function to create a thread.

pthread\_create() takes 4 arguments.

The first argument is a pointer to thread\_id which is set by this function.

The second argument specifies attributes. If the value is NULL, then default attributes shall be used.

The third argument is name of function to be executed for the thread to be created.

The fourth argument is used to pass arguments to the function, myThreadFun.

The pthread\_join() function for threads is the equivalent of wait() for processes. A call to pthread\_join blocks the calling thread until the thread with identifier equal to the first argument terminates.

### **How to Compile the above program:**

To compile a multithreaded program using gcc, we need to link it with the pthreads library. Following is the command used to compile the program.

gcc thread.c -o thread -lpthread ----> for compilation

./thread ----> for execution

OR

gcc -o thread thread.c -lpthread ----> for compilation

./thread -----> for execution

## Example # 02:

Study the format of the code below especially the usage of the bold text.

thread1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
void *print_message_function( void *ptr)
{
    char *message;

    message = (char *) ptr;

    // type casting of pointer to char
    printf("%s \n", message);
}

int main()
{
    pthread_t thread1, thread2;
    char *message1 = "Thread 1";
    char *message2 = "Thread 2";
    int return_value1, return_value2;

    // Create independent threads each of which
    // will execute the same function
```



```

return_value1=pthread_create(&thread1,NULL,
print_message_function, message1);

return_value2=pthread_create(&thread2,NULL,
print_message_function, (void*) message2);

// Wait till threads are complete before main continues.

pthread_join( thread1, NULL);

pthread_join( thread2, NULL);

printf("Thread 1 returns: %d\n",return_value1);
printf("Thread 2 returns: %d\n",return_value2);

exit(0);

return 0;

}

```

Compile the above code. Compilation of multithreaded programs are differently from the normal method. For threads, the `-lpthread` argument is provided as an addition.

When this program runs, there will be a total of 3 threads running `_in_` this process; main, thread1, and thread2. Three additional kernel level threads will be required for servicing these three user-level threads. (Remember, Linux uses 1:1 thread model).

When the thread is created using `pthread_create()`, the main thread proceeds executing with the remainder of instructions. Meanwhile, the newly created thread will complete executing as well. If in case the `main()` thread finishes executing before any of the other threads, the process will exit. Any thread which had not finished will be interrupted before its termination. To avoid this, the `pthread_join()` call can be used. With this, the `main()` thread will wait for successful completion of any thread specified in the join call.

Actually, `pthread_join()` is the opposite of `pthread_create()`. `Pthread_create()` will split our single threaded process into two-threaded process. `Pthread_join()` will join back the two threaded process into a single threaded process.

Read the manual pages and find out the answers to the following:

**Q1:** What is the pthread\_create() and the pthread\_join() calls doing?

**Q2:** In the pthread\_create() call, what are the 4 parameters?

**Q3:** In pthread\_create() call, the 4th parameter is used for passing a pointer to argument of a function. What will we need to do if we want to pass multiple arguments to that function?

**Q4:** In the pthread\_join() call, what are the 2 parameters?

**Q5:** What is the purpose of the return\_value1 and return\_value2 variables? Find out the contents of both these variables using a printf function. What do they contain?

## Passing Multiple Arguments to Thread

Multiple arguments can be passed to a thread through declaring a structure. For example,

```
struct thread_data
{
    int x;
    int y;
    int z;
};
```

Would be our thread\_data structure containing members x, y, and z.

```
struct thread_data Omar;    /* structs to be passed to threads */
```

Will create an instance of thread\_data structure. Its name is given as **Omar**.

```
void print_message_function ( void *ptr )
{
    struct thread_data *my_data;
```

```

    my_data = (thdata *) ptr; /* type cast to a pointer to thdata */
    /* do the work */
    printf("X: %d, Y: %d, Z: %d\n", my_data->x, my_data->y, my_data->z);
    pthread_exit(0);          /* exit */
} /* print_message_function ( void *ptr ) */

```

Is our thread function. Here, we have declared a pointer called my\_data and assign it location of argument passed to thread as input. Since this is a pointer to a structure, we will be using the arrow operator (->) instead of the dot (.) operator to access its members.

```

int main()
{
    pthread_t thread1; /* thread variables */
    struct thread_data Omar; /* structs to be passed to threads */
    Omar.x=1;
    Omar.y=2;
    Omar.z=3;
    pthread_create (&thread1, NULL, (void *) &print_message_function, (void *) &Omar);
    pthread_join(thread1, NULL);
    exit(0);
}

```

Here, we assign the members x, y, and z some values. Instead of sending these individual

members over to the thread, we send the address of the instance Omar.

### Example # 03

struct\_thread.c

```

#include <unistd.h>
#include <sys/types.h> /* Primitive System Data Types */
#include <errno.h> /* Errors */

```

```

#include <stdio.h> /* Input/Output */
#include <stdlib.h> /* General Utilities */
#include <pthread.h> /* POSIX Threads */
#include <string.h> /* String handling */
// prototype for thread routine

void print_message_function ( void *ptr );

/*struct to hold data to be passed to a thread this shows
how multiple data items can be passed to
a thread */

struct thread_data
{
    int x;
    int y;
    int z;
};

int main()
{
    pthread_t thread1; /* thread variables */
    struct thread_data Omar;

    /* initialize data to pass to thread 1 */
    Omar.x=1;
    Omar.y=2;
    Omar.z=3;

    pthread_create(&thread1, NULL, (void *) &print_message_function,
        (void *) &Omar);

```

```

pthread_join(thread1, NULL);

exit(0);

}

void print_message_function ( void *ptr )
{
    struct thread_data *my_data;

    /* type cast to a pointer to thread_data */
    my_data = (struct thread_data *) ptr;

    printf("X: %d, Y: %d, Z: %d\n", my_data->x, my_data->y,
my_data->z);

    pthread_exit(0); /* exit */
}

```

## Thread Termination

There are several ways of terminating threads. Summarised as follows:

- The thread `_returns_` from its starting routine.
- Thread makes a call to `pthread_exit()` call.
- The entire process is terminated due to call to `exec()` or `exit()`.

If the `main()` thread finishes executing before any other thread in the process, all threads will terminate. However, if `main()` exits using a `pthread_exit()` call, all other threads in the process will continue to execute. Thus the `pthread_exit()` call will terminate the main thread but keep on clinging to resources such as process memory space and open file descriptors.

The following code will show both thread creation and thread termination.

## Example # 04

threadTermination.c

```
#include <pthread.h>

#include <stdio.h>

#include <stdlib.h>

void *PrintHello()

{

    printf("Hello World! It's me\n");

    pthread_exit(0);

}

int main()

{

    pthread_t threads[3];

    int rc;

    int t;

    for(t=0; t<3; t++)

    {

        printf("In main: creating thread %d\n", t);

        rc = pthread_create(&threads[t], NULL, PrintHello, NULL);

        if (rc)

        {

            printf("ERROR; return code from pthread_create() is %d\n", rc);

            exit(-1);
```

```

        }

    }

    pthread_exit(NULL);
}

```

## Data Sharing between Threads

### Example # 05

Compile and run the following code:

**threadSharingData.c**

```

#include <pthread.h>

#include <stdlib.h>

#include <unistd.h>

#include <stdio.h>

int myGlobal = 0;

void *threadFunction() {

    int i, j;

    for (i = 0; i<5; i++)

    {

        j = myGlobal;

        j = j+1;

        printf(".");

        fflush(stdout);

        sleep(1);

        myGlobal = j;

    } }

```

```

int main()
{
pthread_t myThread;

int i;

pthread_create(&myThread, NULL, threadFunction, NULL);

for (i = 0; i < 5; i++){

    myGlobal = myGlobal + 1;

    printf("o");

    fflush(stdout);

    sleep(1);

}

pthread_join(myThread, NULL);

printf("\nMy Global Is: %d\n", myGlobal);

exit(0);

}

```

Flush() is used to force a write of all user-space buffered data. Since we have specified stdout, therefore it will write it to standard output. We have two threads here again as well. The main thread has a for loop which is printing the character “o”. The threadFunction thread also has a for loop which is printing the character “.”. You will notice thread-scheduling in action when you see an output of:

### Output:

```
usman@usman-7G-Series:~/os_lab/lab12$ ./threadSharingData
```

```
o.o.o.o.o.
```

```
My Global Is: 6
```



Is everything fine with this output? Think about the myGlobal value ... Should it be 6?

Or should it be 10 (2 For loops each running for 5 iterations)?

Now comment the sleep line (just one) and check the output. It should be 10.

How come a sleep(1) can make such a difference? Here is a little explanation:

We are using the sleep() call to impose a rudimentary form of synchronization between both threads. With sleep, CPU alternates between both threads a total number of 5 times. Each time the myGlobal variable is overwritten by the threadFunction thread. Without sleep, there is just one alternation. The my-Global variable first counts upto 5, and then it counts upto 5 again, totalling 10. But there is no synchronization between threads in this way.

If we want to impose synchronization and at the same time preserve data integrity, we would be needing something more accurate than a simple sleep() call.

## **Synchronization through Mutex**

Every process has a certain portion of code which is called the “critical section” of a process. As an example, this is the area where a

process may:

1. Change shared variables
2. Write to a File
3. Use a shared resource

It would be desirable that if one process is working in this region, then another process should not be allowed to modify the contents of any shared variable within it. If, and only if, that process leaves the critical section, then another process may be allowed access to that shared region. In other words, such a region is “mutually exclusive” to one-and-only-one process at a time.

Ideally, if a process enters this region, it should lock it. Any process trying to access it in the meanwhile will not gain any access because it is locked. Once the process leaves this region, it will un-lock it, rendering it open for anybody else.

We will take this concept of critical-section and mutual exclusion and apply it to our problem in this Section. Here, we apply our lock and unlock mechanism with the help of

Mutexes (taken from Mutually Exclusives). If a thread is currently locked into its critical section, another thread trying to access it will go into sleep mode. Compile and run the code given. Here, the lock is our entry section and the unlock is our exit section.

## Example # 06

mutexSynchronization.c

```
#include <pthread.h>

#include <stdlib.h>

#include <unistd.h>

#include <stdio.h>

int myGlobal = 0;

pthread_mutex_t myMutex;

void *threadFunction()
{
    int i, j;

    for (i = 0; i<5; i++)
    {
        pthread_mutex_lock(&myMutex);

        j = myGlobal;

        j = j+1;

        printf(".");

        fflush(stdout);

        sleep(1);

        myGlobal = j;

        pthread_mutex_unlock(&myMutex);
    }
}
```

```

}

int main()
{
    pthread_t myThread;

    int i;

    pthread_create(&myThread, NULL, threadFunction, NULL);

    for (i = 0; i < 5; i++)
    {
        pthread_mutex_lock(&myMutex);

        myGlobal = myGlobal + 1;

        pthread_mutex_unlock(&myMutex);

        printf("o");

        fflush(stdout);

        sleep(1);
    }

    pthread_join(myThread, NULL);

    printf("\nMy Global Is: %d\n", myGlobal);

    exit(0);
}

```