


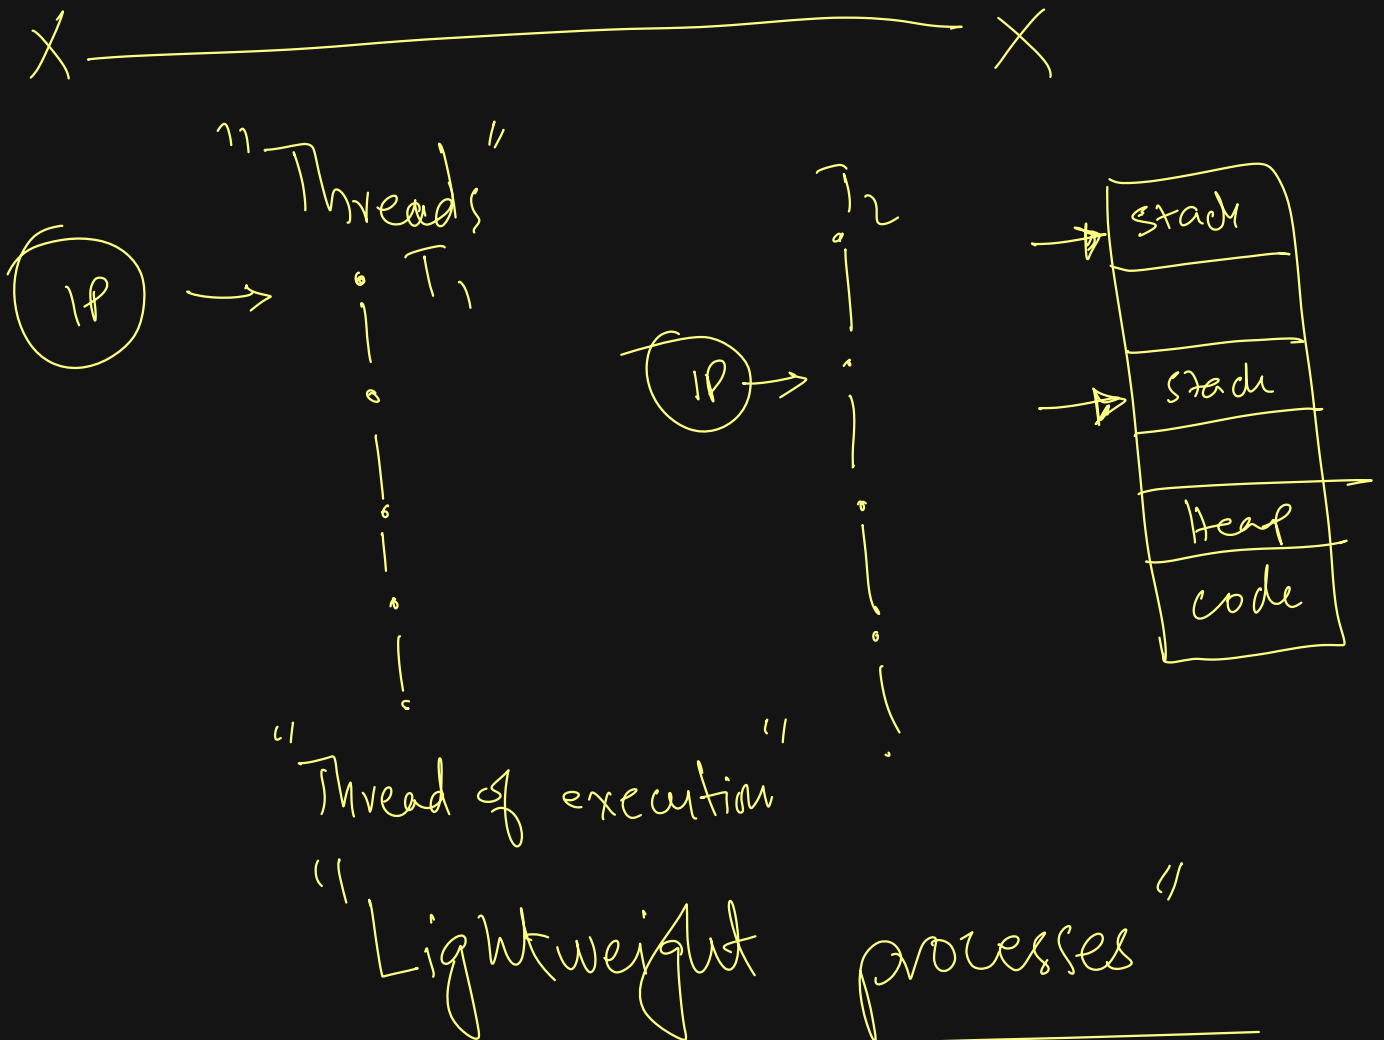
Operating Systems Design

5. Threads

Paul Krzyzanowski
pxk@cs.rutgers.edu



fork is not "heavy weight".



Thread of execution

Single sequence of instructions

- Pointed to by the program counter (PC)
- Executed by the processor

Conventional programming model & OS structure:

- Single threaded
- One process = one thread

Multi-threaded model

A thread is a subset of a process:


- A process contains one or more kernel threads

Share memory and open files

- **BUT:**
separate program counter, registers, and stack
- Shared memory includes the heap and global/static data
- **No memory protection among the threads**

Preemptive multitasking:

- Operating system preempts & schedules threads



Sharing

Threads share:

- Text segment (instructions)
- Data segment (static and global data)
- BSS segment (uninitialized data)
- Open file descriptors
- Signals
- Current working directory
- User and group IDs

Threads do not share:

- Thread ID
- Saved registers, stack pointer, instruction pointer
- Stack (local variables, temporary variables, return addresses)
- Signal mask
- Priority (scheduling information)

Why is this good?

Threads are more efficient

- Much less overhead to create: no need to create new copy of memory space, file descriptors, etc.

Sharing memory is easy (automatic)

- No need to figure out inter-process communication mechanisms

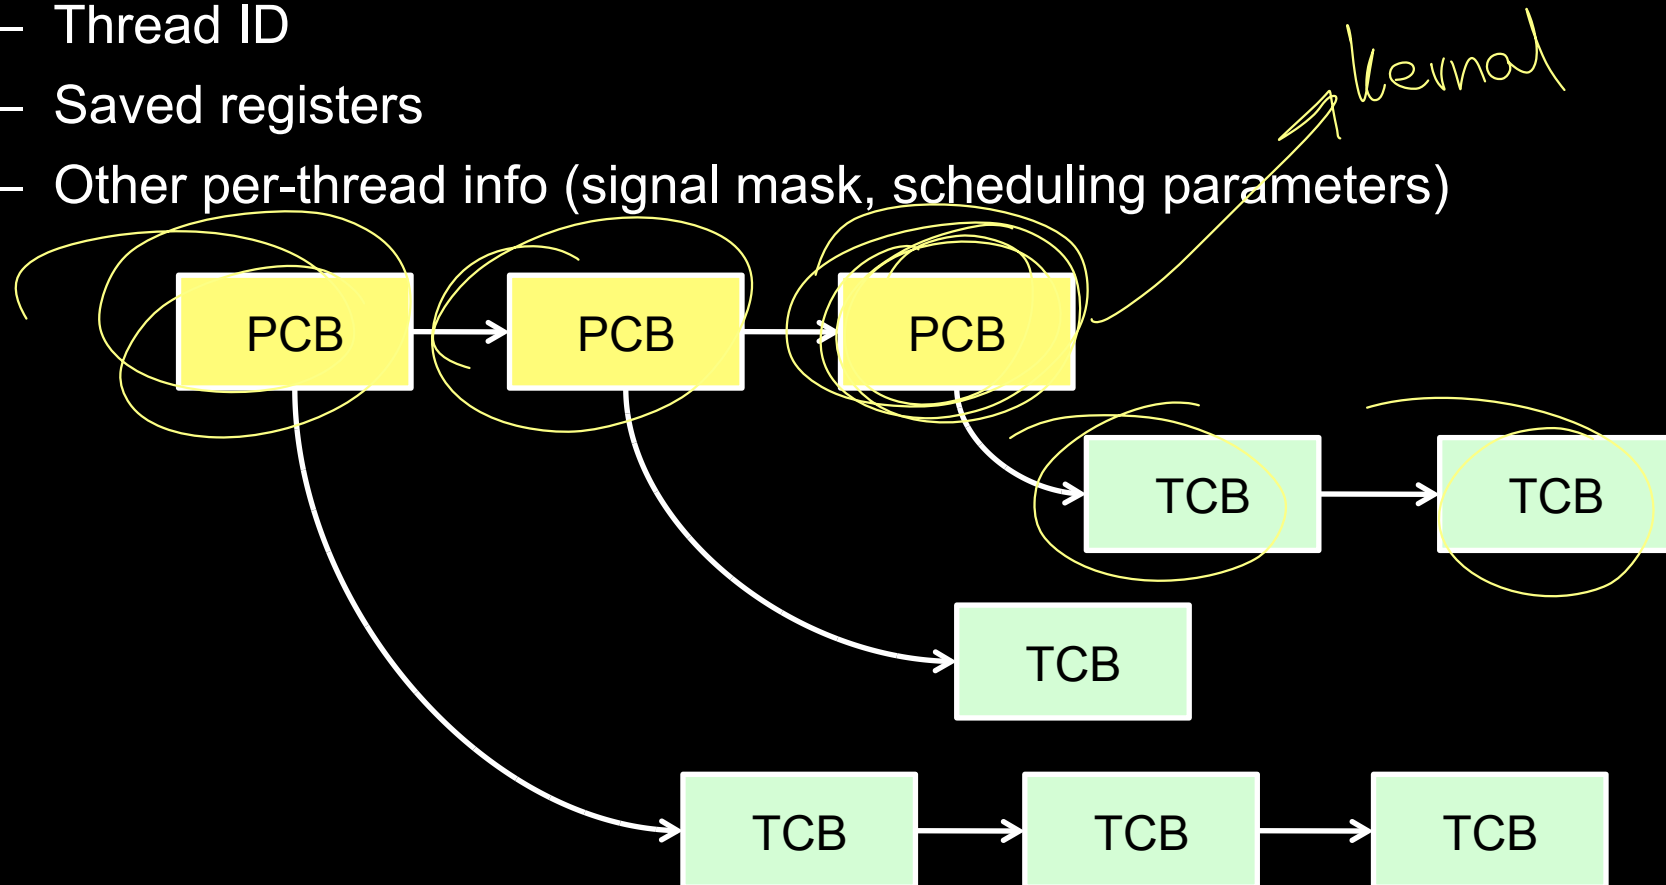
Take advantage of multiple CPUs – just like processes

- Program scales with increasing # of CPUs
- Take advantage of multiple cores

Implementation

Process info (Process Control Block) contains one or more Thread Control Blocks (TCB):

- Thread ID
- Saved registers
- Other per-thread info (signal mask, scheduling parameters)



Scheduling

A threaded-aware operating system scheduler schedules threads, not processes

- A process is just a container for one or more threads

Scheduler has to realize:


- Context switch among threads of different processes is more expensive:
 - Flush cache memory (or have memory with process tags)
 - Flush virtual memory TLB (or have tagged TLB)
 - Replace page table pointer in memory management unit
- Scheduling threads onto a different CPU is more expensive
 - The CPU's cache may have memory used by the thread cached
 - CPU affinity

Process vs. Thread context switch

[A thread switch within the same process is not a context switch]

linux/arch/i386/kernel/process.c:

```
/* Re-load page tables for a new address space */
{
    unsigned long new_cr3 = next->tss.cr3;
    if (new_cr3 != prev->tss.cr3)
        asm volatile("movl %0,%%cr3": : "r" (new_cr3));
}
```



Programming patterns

Single task thread

- Do a specific job and then release the thread

Worker threads

- Specific task for each worker thread
- Dispatch task to the thread that handles it

Thread pools

- Create a pool of threads *a priori*
- Use an existing thread to perform a task; wait if no threads available
- Common model for servers

Kernel-level threads vs. User-level threads

Kernel-level

- Threads supported by operating system
- OS handles scheduling, creation, synchronization

User-level

- Library with code for creation, termination, scheduling
- Kernel sees one execution context: **one process**
- May or may not be preemptive

User-level threads

Advantages

- Low-cost: user level operations that do not require switching to the kernel
- Scheduling algorithms can be replaced easily & custom to app
- Greater portability

Disadvantages

- If a thread is blocked, all threads for the process are blocked
 - Every system call needs an asynchronous counterpart
- Cannot take advantage of multiprocessing

You can have both

User-level thread library on top of multiple kernel threads

1:1 – pure kernel threads only
(1 user thread = 1 kernel thread)

N:1 – pure user threads only
(N user threads on 1 kernel thread/process)

N:M – hybrid threading
(N user threads on M kernel threads)

pthread: POSIX Threads

- POSIX.1c, Threads extensions (IEEE Std 1003.1c-1995)
- Defines API for managing threads
- Linux: native POSIX Thread Library (as of 2.6 kernel)
- Also on Solaris, Mac OS X, NetBSD, FreeBSD
- API library on top of Win32

Using POSIX Threads

Create a thread

```
pthread_t t;  
pthread_create(&t, NULL, func, arg)
```

- Create new thread *t*
- Start executing function *func(arg)*

Join two threads:

```
void *ret_val;  
pthread_join(t, &ret_val);
```

- Wait for thread *t* to terminate (via *return* or *pthread_exit*)

No parent/child relationship!

- Any one thread may wait (join) on another thread


Linux *clone()* system call

- Clone a process, like *fork*, but:
 - Specify function that the child will run (with argument)
 - Child terminates when the function returns
 - Specify location of the stack for the child
 - Specify what's shared:
 - Share memory (otherwise memory writes use new memory)
 - Share open file descriptor table
 - Share the same parent
 - Share root directory, current directory, and permissions mask
 - Share namespace (mount points creating a directory hierarchy)
 - Share signals
 - *And more...*
- Used by pthreads

Threading in hardware

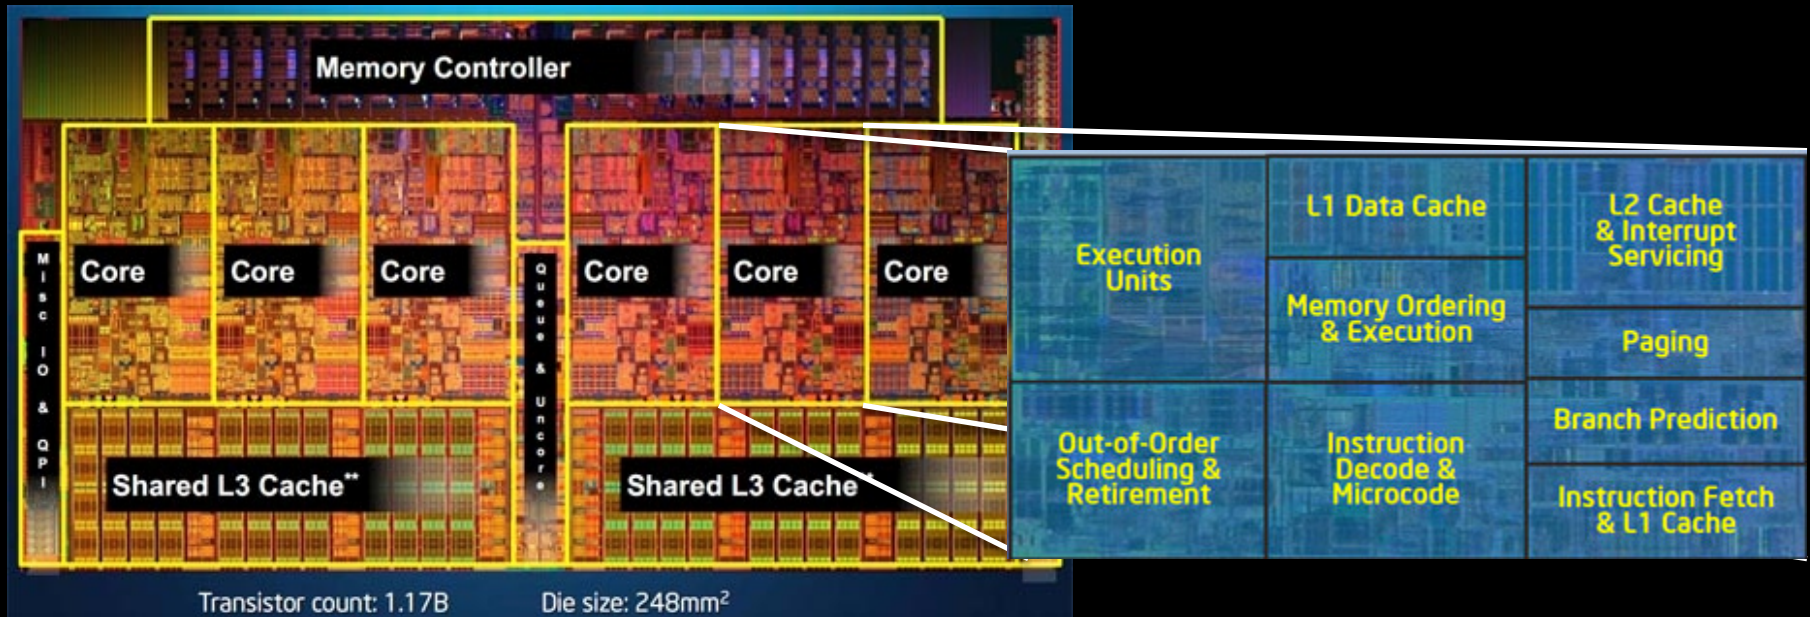
- Hyper-Threading (HT) vs. Multi-core vs. Multi-processor
- One core = One CPU
- Hyper-Threading
 - One physical core *appears* to have multiple processors
 - Looks like multiple CPUs to the OS
 - Multiple threads run but compete for execution unit
 - Events in the pipeline switch between the streams
 - Threads do not have to belong to the same process
 - But the processors share the same cache
 - Performance can degrade if two threads compete for the cache
 - Works well with instruction streams that have large memory latencies

Multi-core architecture



Example CPU

- Intel® Core™ i7-980X Processor Extreme Edition (Sandy Bridge) 3.3 GHz up to 3.6 GHz (Turbo)
- 6 cores; 12 threads
- Per-Core caches:
 - 64 KB L1 cache (32 KB data; 32 KB instruction)
 - 256 KB L2 cache
- 12 MB L3 cache



Stepping on each other

- Threads share the same data
- Mutual exclusion is critical
- Allow a thread be the only one to grab a **critical section**
 - Others who want it go to sleep

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;  
...  
pthread_mutex_lock(&m);  
/* modify shared data */  
pthread_mutex_unlock(&m);
```

The End